


The background features a complex network of red lines connecting green dots, resembling a neural network or a data visualization. This is overlaid on a light blue and white geometric pattern. In the bottom right corner, there is a small, faint number '1'.

Neural Networks and Deep Learning



Outline

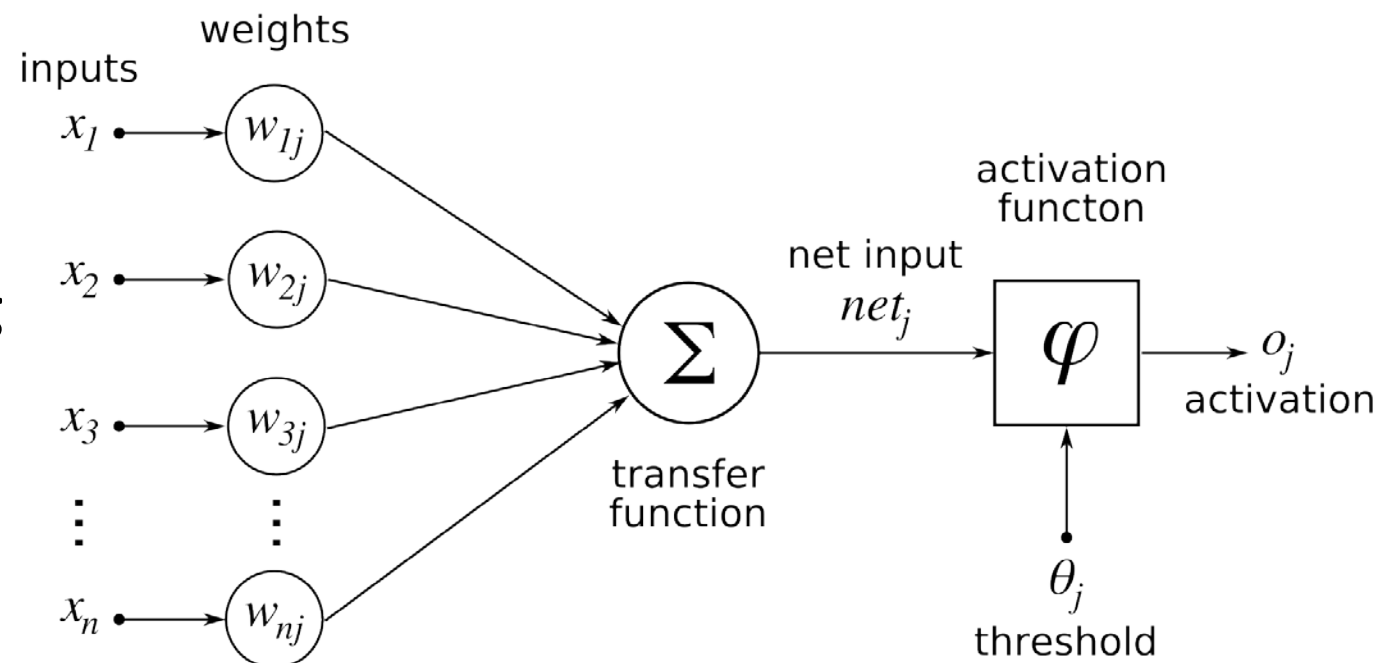
- ❑ Neural Networks
- ❑ Deep Learning: A Short Introduction

The background is a complex, abstract composition. It features a central white banner with a subtle, light gray geometric pattern. Above and below this banner are sections with a dark, reddish-brown background, overlaid with a network of thin, light-colored lines and small, colorful dots (green, blue, yellow). On the left side, there is a vertical strip with a light blue background, containing a grid of small, colorful squares (orange, red, blue, green) and a larger, more complex pattern of dots and lines. The overall aesthetic is modern and technological.

Neural Networks


Neural Network for Classification

- Started by psychologists and neurobiologists to develop and test computational analogues of neurons
- A neural network: A set of connected input/output units where each connection has a **weight** associated with it
- During the learning phase, the **network learns from the training examples by adjusting the weights** so as to be able to predict the correct class label
- Learning without task-specific programming



Artificial Neural Networks as an analogy of Biological Neural Networks

Perceptron: Predecessor of a Neural Network



$$\text{output} = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

- ❑ The perceptron algorithm: Invented in 1957 by Frank Rosenblatt
- ❑ Input: An n -dimensional input vector \mathbf{x} (with n variables)
- ❑ Output: 1 or 0 depending on if the weighted sum passes a threshold
- ❑ Perceptron: A device that makes decisions by weighing up evidence
- ❑ Often written in the vector form, using bias (b) instead of threshold, as

$$\text{output} = \begin{cases} 0 & \text{if } w \cdot x + b \leq 0 \\ 1 & \text{if } w \cdot x + b > 0 \end{cases}$$


Bias: A measure of how easy it is to get the perceptron to output a 1

Sigmoid Neurons

- ❑ A many-layer network of perceptrons can engage in sophisticated decision making
- ❑ Instead of assigning weights of the edges by a person, we can devise *learning algorithms* that can automatically tune the weights and biases of a network of artificial neurons
- ❑ Use sigmoid neuron instead of perceptron: Output is not 0/1 but a *sigmoid function*: $\sigma(w \bullet x + b)$, i.e.,
- ❑ The smoothness of σ means that small changes in the Δw_j weights and in the Δb bias will produce a small change Δ_{output} in the output from the neuron

$$\Delta_{\text{output}} \approx \sum_j \frac{\partial \text{output}}{\partial w_j} \Delta w_j + \frac{\partial \text{output}}{\partial b} \Delta b$$


i.e., Δ_{output} is a *linear function* of the changes Δw_j and Δb



© 2005 Charlie / CC BY-SA 2.5 / <https://goo.gl/PKrbqF>

Sigmoid function: $\sigma(z) \equiv \frac{1}{1 + e^{-z}}$

$$\frac{1}{1 + \exp(-\sum_j w_j x_j - b)}$$



Architecture of (Feed-Forward) Neural Network (NN)

□ Input layer


- The **inputs** to a neural network correspond to the attributes measured for each training instance
- Inputs are fed simultaneously into the units making up the **input layer**

□ Hidden layer(s)

- Inputs are weighted and fed simultaneously to a hidden layer
- The number of hidden layers is arbitrary

□ Output layer

- The weighted outputs of the last hidden layer are input to units making up the **output layer**, which emits the network's prediction



© 2005 Charlie / CC BY-SA 2.5 / <https://goo.gl/PKrbqF>

Neural Network Architecture: Feed-Forward vs. Recurrent

- ❑ **Feed-Forward Neural Network:** Typical neural network architecture
 - ❑ The output from one layer is used as input to the next layer (no loops)
 - ❑ Information is always fed forward, never fed back
 - ❑ From a statistical point of view, networks perform **nonlinear regression**
 - ❑ Given enough hidden units and enough training samples, they can closely approximate any function
- ❑ **Recurrent neural network:** Feedback loops are possible (cascade of neurons firing)
 - ❑ Some neurons fire for some limited duration of time, before becoming quiescent
 - ❑ That firing can stimulate other neurons, which may fire a little while later, also for a limited duration, which causes still more neurons to fire, and so on
 - ❑ Loops do not cause problems since a neuron's output only affects its input at some later time, not instantaneously

Learning with Gradient Descent

- A quadratic cost (objective) function C (or mean square error, MSE)

$$C(w, b) \equiv \frac{1}{2n} \sum_x \|y(x) - a\|^2$$

where w : The collection of all weights in the network; b : All the biases; n : The total # of training inputs; a : The vector of outputs from the network when x is input

- Goal of training a network: Find weights and biases which minimize the cost $C(w, b)$
- For two variables, it means: Choose Δv_1 and Δv_2 to make negative; i.e., the ball is rolling down into the valley:


$$\Delta C \approx \frac{\partial C}{\partial v_1} \Delta v_1 + \frac{\partial C}{\partial v_2} \Delta v_2$$

- The change ΔC in C by a small change in v , Δv :


$$\Delta C \approx \nabla C \cdot \Delta v$$

where ∇C is the gradient vector:

$$\nabla C \equiv \left(\frac{\partial C}{\partial v_1}, \dots, \frac{\partial C}{\partial v_m} \right)$$



Stochastic Gradient Descent




© 2015 Роман Сузи/ CC BY-SA 3.0 /<https://goo.gl/31D7JZ>

- Gradient descent can be viewed as a way of taking small steps in the direction, which does the most to immediately decrease C
- To compute gradient ∇C , we need to compute the gradients ∇C_x separately for each training input, x , and then average them: Slow when the # of training inputs is large
- *Stochastic gradient descent* (SGD): Speed up learning
 - Computing for a small sample of randomly chosen training inputs and *averaging over them*, we can quickly get a good estimate of the true gradient
 - Method: Randomly pick out a small number (**mini-batch**) m of randomly chosen training inputs. Provided the sample size is large enough, we expect that the average value will be roughly equal to the average over all, that is,
$$\nabla C \approx \frac{1}{m} \sum_{j=1}^m \nabla C_{x_j}$$
- Stochastic gradient descent in neural networks:
 - Pick out a randomly chosen minibatch of training inputs and train with them; then pick out another minibatch, until inputs exhausted - complete an *epoch* of training
 - Then we start over with a new training epoch

Backpropagation for Fast Gradient Computation

- ❑ **Backpropagation:** Reset weights on the “front” neural units and this is sometimes done in combination with training, where the correct result is known
- ❑ Iteratively process a set of training instances & compare the network’s prediction with the actual known target value
- ❑ For each training instance, the weights are modified to **minimize the mean squared error** between the network's prediction and the actual target value
- ❑ Modifications are made in the “**backwards**” direction
 - ❑ From the output layer, through each hidden layer back to the first hidden layer, hence “**backpropagation**”
- ❑ Steps
 - ❑ Initialize weights to small random numbers, associated with k
 - ❑ Propagate the inputs forward (by applying activation function)
 - ❑ Backpropagate the error (by updating weights and biases)
 - ❑ Terminating condition (when error is very small, etc.)



© 2005 Charlie / CC BY-SA 2.5 /
<https://goo.gl/PKrbqF>

More on Backpropagation

- ❑ With backpropagation, we distribute the “blame” backward through the network
 - ❑ Each hidden node sending input to the current node is somewhat “responsible” for some portion of the error in each neuron to which it has forward connection
- ❑ Local minima and backpropagation
 - ❑ Backpropagation can be stuck at local minima
 - ❑ But in practice it generally performs well
- ❑ Is backpropagation too slow?
 - ❑ Historically, backpropagation has been considered slow
 - ❑ Recent advances in computer power through parallelism and GPUs (graphics processing units) have reduced time substantially for training neural networks

The background features a complex, abstract design. It includes a grid of small, light-colored plus signs on a pale pinkish-grey field. Overlaid on this are various geometric elements: a network of thin, reddish-brown lines connecting numerous small green dots, and a series of thin, light grey lines forming a triangular mesh. A prominent white, angular shape, resembling a stylized 'V' or a folded piece of paper, serves as a backdrop for the title text. In the bottom right corner, a small, light blue number '1' is visible.

Deep Learning: A Short Introduction


From Neural Networks to Deep Learning

- ❑ Train networks with many layers (vs. shallow nets with just a couple of layers)
 - ❑ More neurons than previous networks
 - ❑ More complex ways to connect layers
 - ❑ Tremendous computing power to train networks
 - ❑ Automatic feature extraction
- ❑ Multiple layers work together to build an improved feature space
 - ❑ Analogy: Signals passing through regions of the visual cortex
 - ❑ Example: For face recognition: Edge → nose → face, layer-by-layer
- ❑ We introduce two popular deep learning frameworks for classification
 - ❑ Convolutional Neural Network (CNN)
 - ❑ Recurrent Neural Network (RNN)




Convolutional Neural Networks: General Architecture

- ❑ Learn high-order features in the data via convolutions
 - ❑ Well suited to object recognition with images (e.g., computer vision)
 - ❑ Build position- and (somewhat) rotation-invariant features from raw image data
- ❑ CNN leverages learnable visual filters and globally shared local features
 - ❑ Specifics: High dimensional, 2D topology of pixels, invariance to translations, etc.
- ❑ High-level general CNN architecture
 - ❑ Input layer
 - ❑ Feature-extraction layers
 - ❑ (Convolution – ReLU – Pool)
 - ❑ Classification layers
- ❑ CNN properties
 - ❑ Local connectivity
 - ❑ Parameter sharing
 - ❑ Subsampling





Convolutional Neural Networks: Local Connectivity

- Convolution: A math operation describing how to merge two sets of information
 - Filter (kernel): Sets of weights in a convolutional layer $(x * k)_{ij} = \sum_{pq} x_{i+p, j+q} k_{r-p, r-q}$
 - The filter is convolved with the input, resulting in a feature (activation) map
- Local Connectivity
 - Receptive fields: Each hidden unit is connected only to a sub-region of the image
 - Manageable number of parameters
 - Efficient computation of pre-activation
 - Spatial arrangements
 - Output Depth: Number of filters
 - Stride: How far our slide filter window will move
 - Zero-padding: Dealing with the border




Convolutional Neural Networks: Parameter Sharing

- Parameter sharing
 - Discrete convolution: Share matrix of parameters across certain units
 - Reduces even more the number of parameters
 - Extract the same feature at every position



Convolutional Neural Networks: Subsampling

- ❑ Pooling layers are commonly inserted between successive convolutional layers
- ❑ Subsampling:
 - ❑ Pooling: Pool hidden units in the same neighborhood
 - ❑ Introduces invariance to local translations
 - ❑ Reduces the number of hidden units in hidden layer



Single depth slice

x	1	1	2	4
	5	6	7	8
	3	2	1	0
	1	2	3	4
	y			

Max pooling: The most common downsampling operation

max pool with 2×2 filters and stride 2

6	8
3	4



Recurrent Neural Networks: General Concepts

- ❑ Modeling the time dimension: By creating cycles in the network (thus “recurrent”)
 - ❑ Adding feedback loops connected to past decisions
 - ❑ Long-term dependencies: Use hidden states to preserve sequential information
 - ❑ Thus, it allows for both parallel and sequential computations
- ❑ Recurrent Neural Networks (RNNs) are trained to generate sequences: Output at each timestamp is based on both the current input and the inputs at all previous timestamps:
$$\mathbf{h}_t = \phi(W\mathbf{x}_t + U\mathbf{h}_{t-1}),$$
 - ❑ Compute a gradient with algorithm BPTT (backpropagation through time)
- ❑ Major obstacles of RNN: Vanishing and exploding gradients
 - ❑ When the gradient becomes too large or too small, it is difficult to model long-range dependencies (10 timestamps or more)
 - ❑ Solution: Use a variant of RNN: LSTM (Long Short-Term Memory) (by Hochreiter and Schmidhuber, 1997)

LSTM: A Variant of Recurrent Neural Network

❑ Critical components of LSTM

- ❑ Memory cells

- ❑ 3 Gates (input, forget, output)

❑ Use gated cells to

- ❑ Write, store, forget information

❑ When both gates are closed

- ❑ The contents of the memory


❑ The gating structure allows information to be retained across many timestamps

- ❑ Also allows gradient to flow across many timestamps

❑ By back-propagating errors and adjusting weights, one can learn what to store, and when to allow reads, writes, and erasures

❑ Applications: Handling sequence and time series data

- ❑ E.g., NLP, video analysis, image captioning, robotics control



Difficulties of Training and Improvements

❑ Challenges

- ❑ Vanishing gradient problem: Saturated units block gradient propagation

 - ❑ Need better optimization than SGD

- ❑ Overfitting: High variance/low bias situation

 - ❑ Need better regularization (than L1, L2 norm)

❑ Many improvements proposed, such as

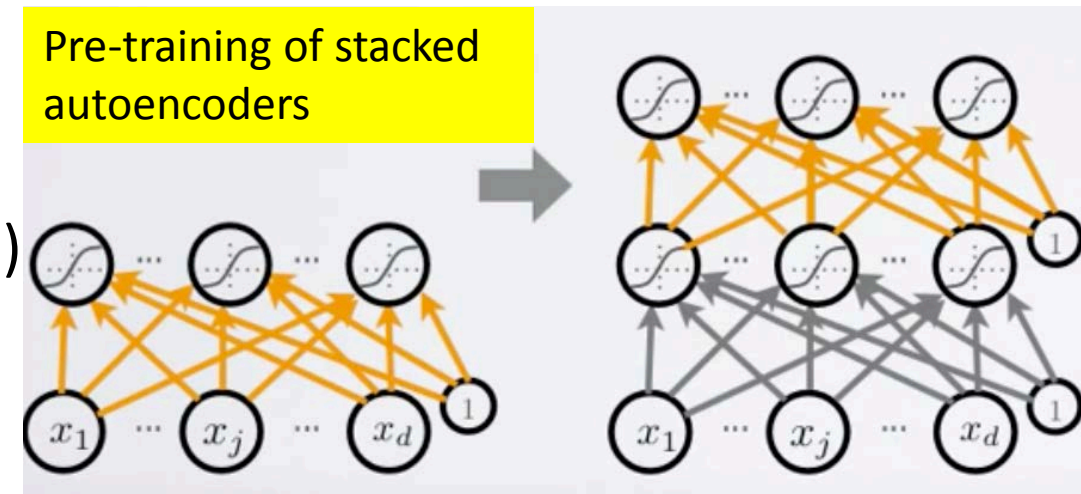
- ❑ Autoencoder

 - ❑ Use unlabeled data in unsupervised learning

 - ❑ Build a compressed representation of the input data

- ❑ Attention: Focusing on specific parts of the input

 - ❑ Taking n arguments (y_1, \dots, y_n) and a context c , it returns a weighted arithmetic mean of the y_i , and the weights are chosen according to y_i 's relevance to c



© 2013 Hugo Larochelle / <https://goo.gl/hNBhRO>



Summary

Summary

- ❑ Neural Networks
- ❑ Deep Learning: A Short Introduction

Recommended Readings

- ❑ Géron, Aurélien. (2017). *Hands-on machine learning with scikit-learn and TensorFlow: Concepts, tools, and techniques to build intelligent systems*. Champaign, IL: O'Reilly.
- ❑ Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. London, UK: The MIT Press.
- ❑ Patterson, J. & Gibson, A. (2017). *Deep learning: A practitioner's approach*. Champaign, IL: O'Reilly.
- ❑ Rashid, Tariq. (2016). *Make your own neural network*. Charleston, SC: CreateSpace.
- ❑ Numerous websites and online tutorials

References

- ❑ Alanf777. (2013). *130316c lc 05 learning curve S-curve or sigmoid.png* [Online image]. Retrieved Feb 19, 2018 from [https://commons.wikimedia.org/wiki/File:130316c lc 05 Learning Curve S-Curve or Sigmoid.png](https://commons.wikimedia.org/wiki/File:130316c_lc_05_Learning_Curve_S-Curve_or_Sigmoid.png)
- ❑ Charlie. (2005). *Neural network.gif* [Online image]. Retrieved Feb 19, 2018 from [https://commons.wikimedia.org/wiki/File:Neural Network.gif](https://commons.wikimedia.org/wiki/File:Neural_Network.gif)
- ❑ Chrislb. (2005). *ArtificialNeuronModel english.png* [Online image]. Retrieved Feb 19, 2018 from https://commons.wikimedia.org/wiki/File:ArtificialNeuronModel_english.png
- ❑ Karpathy, A. & Johnson, J. (2015). *CS231n convolutional neural networks for visual recognition* [Online images]. Retrieved Feb 19, 2018 from <http://cs231n.github.io/convolutional-networks/>
- ❑ Larochelle, H. (2013). Neural networks [7.2] : Deep learning - difficulty of training [Online video]. Retrieved Feb 19, 2018 from https://www.youtube.com/watch?v=YoiUIN_77LU&list=PL6Xpj9I5qXYEcOhn7TqghAJ6NAPrNmUBH&index=52
- ❑ Роман Сузи. (2015). *Gradient descent method.png* [Online image]. Retrieved Feb 19, 2018 from [https://commons.wikimedia.org/wiki/File:Gradient descent method.png](https://commons.wikimedia.org/wiki/File:Gradient_descent_method.png)
- ❑ Skymind. (2017). *Long short-term memory units (LSTMs)* [Online image]. Retrieved Feb 19, 2018 from <https://deeplearning4j.org/lstm.html>
- ❑ All other multimedia elements belong to © 2018 University of Illinois Board of Trustees.