

# CS412 Office Hours

April 5, 2019

# Decision Tree Induction: Algorithm

---

## □ Basic algorithm

- Tree is constructed in a **top-down, recursive, divide-and-conquer manner**
- At start, all the training examples are at the root
- Examples are partitioned recursively based on selected attributes
- On each node, attributes are selected based on the training examples on that node, and a heuristic or statistical measure (e.g., **information gain**)

## □ Conditions for stopping partitioning

- All samples for a given node belong to the same class
- There are no remaining attributes for further partitioning
- There are no samples left

Key issue: How to select the best attribute to partition on?

## □ Prediction

- **Majority voting** is employed for classifying the leaf

# Three Attribute Selection Measures

---

❑ **Information gain:**  $Gain(A) = Info(D) - Info_A(D)$

- ❑ difference between entropy of the response variable  $H(Y)$  and the conditional entropy  $H(Y|A)$
- ❑ biased towards multivalued attributes

❑ **Gain ratio:**  $SplitInfo_A(D) = -\sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2\left(\frac{|D_j|}{|D|}\right)$

- ❑ normalized information gain
- ❑ tends to prefer unbalanced splits in which one partition is much smaller than the others

❑ **Gini index:**  $\Delta gini(A) = gini(D) - gini_A(D)$

- ❑ is biased to multivalued attributes; has difficulty when # of classes is large
- ❑ tends to favor tests that result in equal-sized partitions and purity in both partitions

# Information Gain: An Attribute Selection Measure

---

- ❑ Select the attribute with the highest information gain (used in typical decision tree induction algorithm: ID3)
- ❑ Let  $p_i$  be the probability that an arbitrary tuple in  $D$  belongs to class  $C_i$ , estimated by  $|C_{i,D}|/|D|$
- ❑ Expected information (entropy) needed to classify a tuple in  $D$ :

$$Info(D) = -\sum_{i=1}^m p_i \log_2(p_i)$$

- ❑ Information needed (after using  $A$  to split  $D$  into  $v$  partitions) to classify  $D$ :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

- ❑ Information gained by branching on attribute  $A$

$$Gain(A) = Info(D) - Info_A(D)$$

# Example: Attribute Selection with Information Gain

□ Class P: buys\_computer = “yes”

□ Class N: buys\_computer = “no”

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

age	p <sub>i</sub>	n <sub>i</sub>	I(p <sub>i</sub> , n <sub>i</sub> )
<=30	2	3	0.971
31...40	4	0	0
>40	3	2	0.971

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

$\frac{5}{14} I(2,3)$  means “age <=30” has 5 out of 14 samples, with 2 yes’es and 3 no’s.

Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly, we can get

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit\_rating) = 0.048$$

# Gain Ratio: A Refined Measure for Attribute Selection

---

- ❑ Information gain measure is biased towards attributes with a large number of values
- ❑ Gain ratio: Overcomes the problem (as a normalization to information gain)

$$SplitInfo_A(D) = -\sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2\left(\frac{|D_j|}{|D|}\right)$$

- ❑  $GainRatio(A) = Gain(A)/SplitInfo(A)$
- ❑ The attribute with the maximum gain ratio is selected as the splitting attribute
- ❑ Gain ratio is used in a popular algorithm C4.5 (a successor of ID3) by R. Quinlan
- ❑ Example
  - ❑  $SplitInfo_{income}(D) = -\frac{4}{14} \log_2 \frac{4}{14} - \frac{6}{14} \log_2 \frac{6}{14} - \frac{4}{14} \log_2 \frac{4}{14} = 1.557$
  - ❑  $GainRatio(income) = 0.029/1.557 = 0.019$

# Gini Index

---

- ❑ Gini index: Used in CART, and also in IBM IntelligentMiner
- ❑ If a data set  $D$  contains examples from  $n$  classes, gini index,  $gini(D)$  is defined as
  - ❑  $gini(D) = 1 - \sum_{j=1}^n p_j^2$ 
    - ❑  $p_j$  is the relative frequency of class  $j$  in  $D$
- ❑ If a data set  $D$  is split on  $A$  into two subsets  $D_1$  and  $D_2$ , the  $gini$  index  $gini(D)$  is defined as
  - ❑  $gini_A(D) = \frac{|D_1|}{|D|} gini(D_1) + \frac{|D_2|}{|D|} gini(D_2)$
- ❑ Reduction in Impurity:
  - ❑  $\Delta gini(A) = gini(D) - gini_A(D)$
- ❑ The attribute which provides the smallest  $gini_{split}(D)$  (or the largest reduction in impurity) is chosen to split the node (*need to enumerate all the possible splitting points for each attribute*)

# Computation of Gini Index

- Example: D has 9 tuples in  $\text{buys\_computer} = \text{"yes"}$  and 5 in  $\text{"no"}$

$$\text{gini}(D) = 1 - \left(\frac{9}{14}\right)^2 - \left(\frac{5}{14}\right)^2 = 0.459$$

- Suppose the attribute  $\text{income}$  partitions D into 10 in  $D_1: \{\text{low, medium}\}$  and 4 in  $D_2$

$$\begin{aligned} \text{gini}_{\text{income} \in \{\text{low, medium}\}}(D) &= \frac{10}{14} \text{gini}(D_1) + \frac{4}{14} \text{gini}(D_2) \\ &= \frac{10}{14} \left( 1 - \left(\frac{7}{10}\right)^2 - \left(\frac{3}{10}\right)^2 \right) + \frac{4}{14} \left( 1 - \left(\frac{2}{4}\right)^2 - \left(\frac{2}{4}\right)^2 \right) = 0.443 \\ &= \text{Gini}_{\text{income} \in \{\text{high}\}}(D) \end{aligned}$$

- $\text{Gini}_{\{\text{low, high}\}}$  is 0.458;  $\text{Gini}_{\{\text{medium, high}\}}$  is 0.450

- Thus, split on  $\text{income} \in \{\text{low, medium}\}$  (i.e., also  $\{\text{high}\}$ ) has the lowest Gini index

age	income	student	credit rating	buys computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no




# How to Handle Continuous-Valued Attributes?

---

- Method 1: Discretize continuous values and treat them as categorical values
  - E.g., age: < 20, 20..30, 30..40, 40..50, > 50
- Method 2: Determine the **best split point** for continuous-valued attribute A
  - Sort the value A in increasing order:, e.g., 15, 18, 21, 22, 24, 25, 29, 31, ...
  - *Possible split point*: the midpoint between *each pair of adjacent values*
    - $(a_i + a_{i+1})/2$  is the midpoint between the values of  $a_i$  and  $a_{i+1}$
    - e.g.,  $(15+18)/2 = 16.5, 19.5, 21.5, 23, 24.5, 27, 30, \dots$
  - The point with the *maximum information gain* for A is selected as the **split-point** for A
- Split: Based on split point P
  - The set of tuples in D satisfying  $A \leq P$  vs. those with  $A > P$

# Tree Pruning

- ❑ Two approaches to avoid overfitting
  - ❑ Prepruning: *Halt tree construction early*—do not split a node if this would result in the goodness measure falling below a threshold
    - ❑ Difficult to choose an appropriate threshold
    - ❑ You can experiment with different thresholds for the programming assignment
  - ❑ Postpruning: *Remove branches* from a “fully grown” tree—get a sequence of progressively pruned trees
    - ❑ Use a set of data different from the training data to decide which is the “best pruned tree”



# Using Decision Trees for Large Datasets: RainForest

---

- Our goal is to reduce the number of times we read/write from the disk where the database resides.
- RainForest decouples the “counting” process and the split criteria computation process.
- Database access only happens during the construction of the AVC-groups
- After we get the AVC-groups, we do not need to refer to the original data points anymore

# RainForest: A Scalable Classification Framework

- The criteria that determine the quality of the tree can be computed separately
  - Builds an AVC-list: **AVC (Attribute, Value, Class\_label)**
- **AVC-set** (of an attribute  $X$ )
  - Projection of training dataset onto the attribute  $X$  and class label where counts of individual class label are aggregated
- **AVC-group** (of a node  $n$ )
  - Set of AVC-sets of all predictor attributes at node  $n$
- Then read the data again & do it similarly to generate the next level of the tree

age	income	student	credit rating	comp
<=30	high	no	fair	no
<=30	high	no	excellent	no
31... 40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31... 40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31... 40	medium	no	excellent	yes
31... 40	high	yes	fair	yes
>40	medium	no	excellent	no

The Training Data

AVC-set on Age

Age	Buy_Computer	
	yes	no
<=30	2	3
31..40	4	0
>40	3	2

AVC-set on Income

income	Buy_Computer	
	yes	no
high	2	2
medium	4	2
low	3	1

AVC-set on Student

student	Buy_Computer	
	yes	no
yes	6	1
no	3	4

AVC-set on Credit\_Rating

Credit rating	Buy_Computer	
	yes	no
fair	6	2
excellent	3	3

Its AVC Sets